

# CSE 5031 Operating Systems

## 2020/21 Fall Term

**Project:** 2 – Part 2  
**Topic:** Programming ISAM with Low-level I/O API  
**Date:** 02 - 09.11.2020

### Objectives:

- Implementing Indexed Sequential File Access Method
- Using low-level I/O **GNU C Library API**

### References:

- The GNU C Library Reference Manual (<http://www.gnu.org/software/libc/manual/pdf/libc.pdf> )
- Linux System Programming 2d ed., Robert Love, O'Reilly 2013  
(<http://pdf-ebooks-for-free.blogspot.com.tr/2015/01/oreilly-linux-system-programming.html>)

## Section A. Linux File Concept and I/O APIs

### A.1 UNIX/Linux File Concept

**UNIX** abstracted the **files** stored on any magnetic storage as an **array of bytes** - characters-. **UNIX** extended **file** abstraction over the years to cover:

- ✓ data sets stored on **non-magnetic medium** (various, I/O devices, virtual terminals, inter-process communication channels); and
- ✓ almost all **system entities** that generate or store data, for instance (directories, processes, memory etc.).

To avoid confusions **UNIX/Linux** refers to **files** stored on any magnetic storage as **ordinary files**.

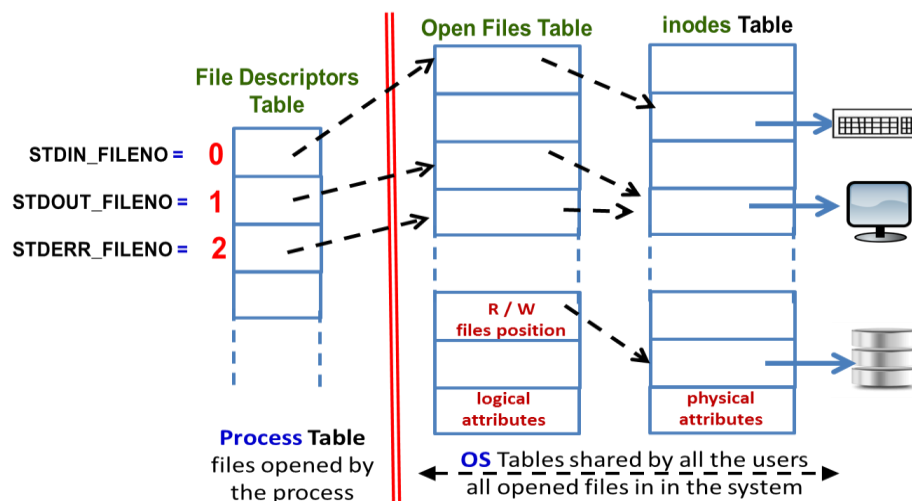
### A.2 GNU UNIX/Linux File Access API

**GNU C library** provides two **APIs** to handle **Linux** files:

- ✓ **low level I/O interface**; and
- ✓ **stream I/O interface**.

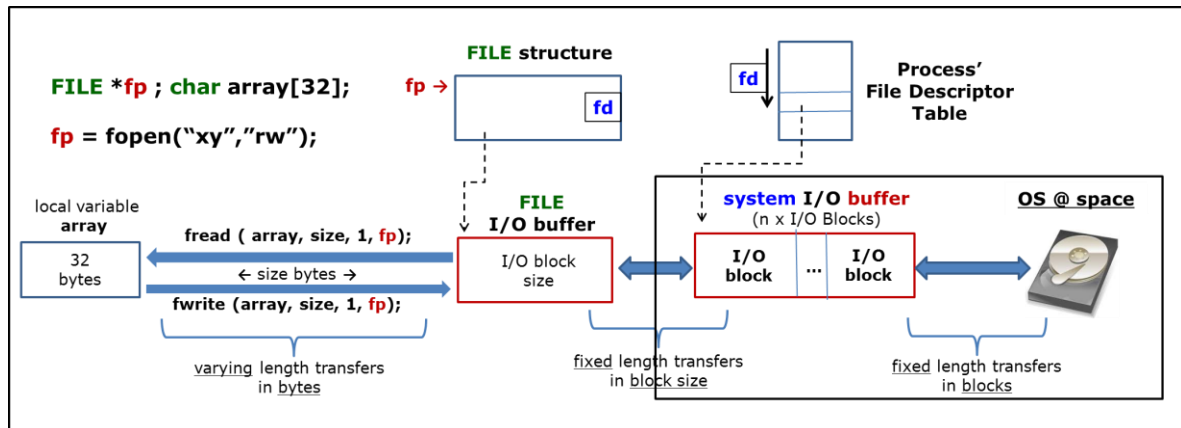
These **APIs** represent the **connection** between a **C** program and **the file entity** using respectively:

- ✓ **File descriptors**, integer ordinals e.g. 0,1,2, ..n referring to an open file in the **File Descriptor Table**, that is stored in the user **virtual address space** in the **OS Kernel zone**; and
- ✓ **FILE Pointers**, referring to a **streams** data structure that is located in the in the user **virtual address space**; that provide a higher-level interface, layered on **top** of the **low-level I/O** mechanism and refer to it using **file descriptor** as shown on the figure at the next page.



Both **API** can also represent a connection to a device (such as a terminal), or a pipe or socket for communicating with another process, as well as to logical system entities such as process and file descriptor tables. But,

- ✓ to perform **control operations** that are specific to a particular kind of device, **low-level I/O API** should be used; it provides only simple functions for transferring **blocks of characters**;
- ✓ the main advantage of using the **stream I/O** interface is its richer and more powerful **I/O functions** set e.g. powerful formatted input & output functions (printf and scanf) as well as functions for character- and line-oriented input & output.
- ✓ **stream I/O** interface has **better performance** in transferring large number of **very small** size I/O data chunks versus low **low-level I/O API** (try to explain why?).



### A.3 Duality of Low-level and Stream I/O APIs

Since **streams** are implemented over **low-level I/O** mechanism, you can extract the **file descriptor** from a **stream** by calling the **"fileno"** function and perform **low-level operations** directly on the file descriptor. You can also initially open a connection as a **file descriptor** by calling the **"fdopen"** function and then make a **stream** associated with that file descriptor (refer to **GNU C Library Reference Manual** section 13.4).

In general, you should stick with using only **one I/O API**, unless there is some specific operation you want to do with the other. If you are concerned about portability of your programs to systems other than **GNU**, **low-level I/O API** be aware that file descriptors **are not as portable** as streams. You can expect any system running **ISO C** to support **streams**, but **non-GNU** systems may not support **low-level IO API** at all or may only implement a subset of it.

### A.4 File Position

**File position** is a **pointer** to the current byte to be read from or written to a file in the byte array abstraction. It is an attribute of an **opened file**. On **GNU** systems, **file position** is represented with an **integer** which counts the **byte offset** – the number of bytes – from the beginning of the file.

When a file is opened for reading or writing, its **file position** is set to the **beginning** of the file, at the **byte offset 0**. Each time a byte is read or written the file position is incremented. In other words, **file access mode** is **sequential**.

Alternatively, an existing file may be opened with the **"append"** attribute to add new records at the end of the file. Such an **open** sets the **file position** to the **end of the file**; if the file size is "**n**" bytes, the offset is set to "**n**" since file offset count starts at 0.

### A.5 Random/Direct Access Modes

**Ordinary files** (data files in UNIX/Linux terminology) permit read or write operations **at any position** within the file. The **file position** may be set:

- ✓ when the file is opened; and successive reads or writes increments the **file position** by the amount of data transferred is or out;
- ✓ to any location using the **fseek** function on a stream (Section 12.18 File Positioning), or the **lseek** function on a file descriptor (Section 13.2 Input and Output Primitives); read / write operations proceed **sequentially** from the **new position**.

This type of **file access mode** is called **sequential** whereas the second **random** or **direct access**.

## Section B. Indexed File Organization

### B.1 Random / Direct File Access the Rationale

Searching a record in a file **sequentially** yields in poor performance when the file contains a **large number** of records. The program must read all the records **one by one** until targeted record is reached, i.e. one or several of its fields are identified by the record search criteria.

Given a file hat contains “n” records, a **sequential** search for a given record requires the reading of “n/2” records **in average**. The process yields in **long access times** and generates a **large number of I/O operations**.

The alternative involves:

- ✓ the acquisition of the record's **file position** (i.e. by looking up in a list of “record key-file position” pairs);
- ✓ setting the **file position** to the position of targeted record (an I/O operation that does not involve a file access but just an update of its data structure in main memory); and
- ✓ reading/writing the record at the new offset (1 read/write).

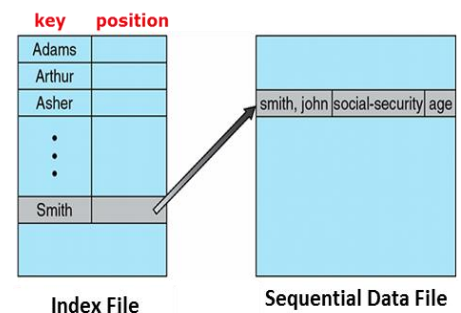
The performance gain achieved by **direct/random** access method is **1 access** versus **n/2 operations**.

### B.2 Indexed File Organization

An effective **direct/random** file access method that enhances the search performance to sequential files is **ISAM**, the *Indexed Sequential Access Method*. **ISAM** uses an **index table** that matches the **key field** of a record with its **file position**. The figure here after depicts a simplified **ISAM** implementation that maps the **family name** of a given customer to its **file position**.

To retrieve a given record using **ISAM** a search program should:

- ✓ acquire the **key** of the record to access;
- ✓ search the **Index Table** in memory for the record **key**;  
if successful:
  - set **file position** to the corresponding record offset extracted from the table using “**lseek**” or “**fseek**” primitive;
  - access the record if positioning is successful.



**ISAM** requires:

- ✓ the generation of a sorted index table associated with the file;
- ✓ reading of the entire “**Index File**” in memory, in the **Index Table**, prior to any access;
- ✓ sorting eventually the data file in the primary key order, in case the **record key is not unique** e.g. several records with the same family name exist; in such case the **Index Table** will be used to access the first record with the same record key and the rest will be browsed in sequence.

### B.3 Creating an ISAM File

In this project you will implement a the **ISAM** file organization for the “**etc/passwd**” file which contains a **unique** record for **each account**. You will build the **index file** using the “**prj2-mkidx.c**”, that:

- ✓ reads the “**passwd**” file, the copy of the “**etc/passwd**” file stored in the current working directory;
- ✓ creates it's **Index Table**, using the first field (the **account name**) as the record key;
- ✓ sorts the **Index Table** in **ascending order** of the record key field;
- ✓ stores the **Index Table** as the “**passwd.idx**” file in the current working directory.

The “**prj2-mkidx.c**” program is stored at the course web site in the folder “**Resources/RefCprograms**” under “**CSE5031-Home**” tab. It uses:

- ✓ **stream I/O API** to read varying length “**passwd**” file records; and
- ✓ **low-level I/O API** to create “**passwd.idx**” file with proper file attributes (ownership and access rights).

Analyze “**prj2-mkidx.c**” taking into consideration the following implementation details:

- ✓ “**passwd**” file is a byte array, containing a **variable length** record per account. Each record consists of **7** varying length fields separated by the “:” character and terminated by the new line (“**\n**”) marker.

```
root:x:0:0:root:/root:/bin/bash\nbin:x:1:1:bin:/bin:/sbin/nologin\ndaemon:x:.....
```

↑ offset 0

↑ offset 32

↑ offset 64

Note that, the record of the account **root** starts at the byte position **0**; the record of **bin** is at offset **32**; and **daemon** account starts at the offset **65** etc.

- ✓ As the fields of “**/etc/passwd**” are variable length, the following assumptions are made on the **max.** lengths:
  - max. lengths for the file name (**NAME\_MAX**), and the path (**PATH\_MAX**) are defined in the system file “**/usr/include/linux/limits.h**”; but they are not used in “**prj2-mkidx.c**”;
  - max. length of the user/account name is defined as **32 byte** in the manual for the “**useradd**” command;
  - max. record length of “**/etc/passwd**” is assumed to be less than **1023** bytes.

## B.4 Creating Project’s ISAM File Index

Perform the following to create the index file “**passwd.idx**”:

- ✓ logon as “**sysadmin**”;
- ✓ create the folder “**prj2**” in your home directory; and set it as your working directory;
- ✓ copy the “**/etc/passwd**” file;
- ✓ copy “**prj2-mkidx.c**” file posted at the course web site;
- ✓ make sure that the owner of “**prj2-mkidx.c**” file is “**sysadmin**”, if not change it using the “**chown**” command;
- ✓ make sure that “**prj2-mkidx.c**” file has **read** and **write** access rights set for the owner; if not change them using the “**chmod**” command;
- ✓ compile “**prj2-mkidx.c**” with gcc, and name it “**mkidx**”;
- ✓ run “**mkidx**” and enter the number of records “**passwd**” contains;
- ✓ make sure that “**passwd.idx**” has been created with correct content and has proper access control rights.

# Section C. Developing the ISAM Query Program

## C.1 Query Program Accessing “passwd” with ISAM

Write a query program in **C** that retrieves a series of selected records from the “**passwd**” file with the **ISAM** method using the Index File “**passwd.idx**” you have created in section B.4.

The query program should:

- ✓ load the file index “**passwd.idx**” in a dynamically allocated **Index Table** (note that you may derive the size of the Index Table from the size of the Index File using the “**stat**” function);
- ✓ read from **standard input** an account name, until an **end of file** is entered (ctrl+del keystrokes for the VM);
- ✓ retrieve the corresponding record from using **ISAM** method;
- ✓ display the **home directory** and **login shell** of the account retrieved if the search is successful.

The query program:

- ✓ must be written using low level I/O API to read the index “**passwd.idx**”; and to access the “**passwd**” file;
- ✓ may use **Stream I/O API** for “**stdin**” and “**stdout**” files.

- ✓ **Chapter 2 “FILE I/O”** of the “**Linux System Programming**” cited under **References** contains authoritative **C** programming examples with the **low-level I/O API**. You are strongly advised to refer to this programming resource, instead of wasting your valuable time in “**fishing junk**” over the Internet.

## C.2 Project Report

- i) Run your query program with several accounts and store the results in an output text file.
- ii) If your **program** is **operational**,
  - add a comment line consisting of your name and student-id;
  - store its **source code** and the **results** files in the “Prj2-Part1” folder, located at the course web site under the tab **CSE5031 - OS Section -X/Assignment**; where “X” stands for (1,2,3,4) the laboratory session group you are registered in.

### Warning

You are encouraged to discuss the implementation procedures and general concepts behind the projects with your fellow students. However, **plagiarism is strictly forbidden!** Submitted report should be the result of **your personal work!**

Be advised that you are **accountable** of your submission not only for this project, but also for the mid-term, and final examinations. Your project grade may be reevaluated retrospectively, had you fail to answer correctly the same or a similar examination questions that you have solved with success in your submissions.